

## Cloud Databases and Data Management

---

This experiment covers setting up and managing cloud-based relational databases using PostgreSQL, Docker, and LocalStack. It provides hands-on experience with database management in a simulated cloud environment. The setup replicates real-world cloud database operations using containerized solutions.

## Cloud Databases and Data Management

### Overview

This experiment covers setting up and managing cloud-based relational databases using PostgreSQL, Docker, and LocalStack. It provides hands-on experience with database management in a **simulated cloud environment**, replicating real-world cloud database operations using containerized solutions.

### What are Databases?

Databases store, organize, and manage data efficiently. They are crucial in applications ranging from websites to large-scale enterprise systems.



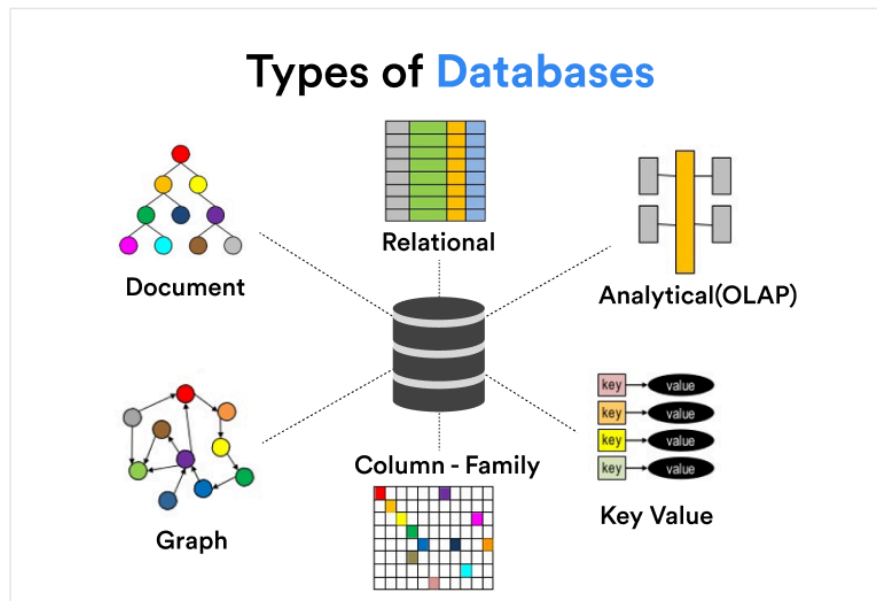
- **A database is an integrated collection of logically related records or files consolidated into a common pool that provides data for one or more multiple uses.**
- **One way of classifying databases involves the type of content, for example: bibliographic, full-text, numeric, image. Other classification methods start from examining [database models](#) or database architectures: see below.**

### Types of Databases

Databases are broadly categorized into:

- **Relational Databases (SQL-based)** – Structured data stored in tables (e.g., PostgreSQL, MySQL).
- **NoSQL Databases** – Flexible schema for handling unstructured data (e.g., MongoDB, Cassandra).


- **Cloud Databases** – Managed services with scalable storage (e.g., AWS RDS, Google Cloud Firestore).



## PostgreSQL Features

PostgreSQL is a powerful, open-source relational database with features like:

- **ACID Compliance** – Ensures data integrity.
- **Extensibility** – Supports custom functions and data types.
- **Scalability** – Handles large volumes of data.



## Real-World Application & Case Study

### Use Case: Financial Data Management

A leading bank implemented PostgreSQL on the cloud to handle transaction processing, fraud detection, and real-time analytics. With **Dockerized deployments**, they achieved high availability, ensuring **99.9% uptime** and enhanced security.

This experiment provides insights into deploying such systems using **Docker** and **LocalStack** to simulate real-world cloud database management.

## Database Operations with Postgres

---

### 1. Creating an RDS Instance Using LocalStack

#### Command:

```
aws rds create-db-instance --db-instance-identifier mydb \  
  --db-instance-class db.t3.micro \  
  --engine mysql \  
  --master-username admin \  
  --master-user-password password \  
  --allocated-storage 20 \  
  --endpoint-url=http://localhost:4566
```

#### Error Output:

```
Could not connect to the endpoint URL: "http://localhost:4566/"
```

```
An error occurred (InternalFailure) when calling the CreateDBInstance operation:  
API for service 'rds' not yet implemented or pro feature - please check  
https://docs.localstack.cloud/references/coverage/ for further information
```

#### Explanation:

- The command attempts to create an **RDS instance** in **LocalStack**.
- `--endpoint-url=http://localhost:4566` → Uses LocalStack instead of AWS.
- **Errors indicate:**
  - LocalStack is either not running or misconfigured.
  - RDS API might not be fully implemented in the **free** version of LocalStack.

#### Output Breakdown:

- **Could not connect to the endpoint URL** → LocalStack might not be running or accessible.
- **InternalFailure error** → The RDS API might require LocalStack Pro for full functionality.
- **Possible Fixes:**
  - Ensure LocalStack is running:

```
docker run --rm -d --name localstack_main -p 4566:4566 localstack/localstack
```

- Check service coverage:  
[LocalStack RDS Coverage](#)

## 2. Starting a PostgreSQL Container

### Command:

```
docker start my-postgres
```

```
docker start postgres
```

### Error Output:

```
Error response from daemon: No such container: my-postgres
Error: failed to start containers: my-postgres
```

```
Error response from daemon: No such container: postgres
Error: failed to start containers: postgres
```

### Explanation:

- The **containers do not exist** under the specified names.
- Verify running containers with:

```
docker ps -a
```

- If needed, create a new container:

```
docker run --name my-postgres -e POSTGRES_USER=admin \
  -e POSTGRES_PASSWORD=password -e POSTGRES_DB=mydb \
  -p 5432:5432 -d postgres:15
```

### Output Breakdown:

- **No such container: my-postgres** → The container was never created or was removed.
- **failed to start containers** → The container name does not match any existing instances.
- **Possible Fixes:**
  - Check existing containers: `docker ps -a`
  - Create and start a new PostgreSQL container using `docker run` (above).

### 3. Listing Available Docker Images

Command:

```
docker images
```

Output:

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
-----	-----	-----	-----	-----
my-flask-app	latest	f5feae0ac7a4	6 hours ago	139MB
flask-app	latest	ae4054c49614	7 hours ago	139MB
hackvortex-backend	latest	14e63c26d40b	21 hours ago	1.05GB
postgres	15	e45d3f5ec589	7 days ago	430MB
localstack/localstack	latest	b686f3948f42	6 weeks ago	1.18GB
python	3.9	9f98746e2033	3 months ago	999MB
nginx	latest	b52e0b094bc0	4 weeks ago	192MB

Explanation:

- Displays **available images** in the local Docker environment.
- PostgreSQL ( `postgres:15` ) is available.
- LocalStack ( `localstack/localstack` ) is present but needs verification ( `docker ps -a` ).

Output Breakdown:

- **postgres:15 is listed** → The image exists but the container may not be running.
- **localstack/localstack exists** → LocalStack is installed but may need to be started.
- **Possible Fixes:**
  - Start PostgreSQL if not running:

```
docker run --name my-postgres -e POSTGRES_USER=admin \  
-e POSTGRES_PASSWORD=password -e POSTGRES_DB=mydb \  
-p 5432:5432 -d postgres:15
```

- Ensure LocalStack is running:

```
docker start localstack_main
```

### 4. Starting a PostgreSQL Container

Command:

```
C:\Users\rawat>docker start postgres
```

Error Output:

```
Error response from daemon: No such container: postgres
Error: failed to start containers: postgres
```

5. Listing All Containers

Command:

```
C:\Users\rawat>docker ps -a
```

Output:

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
a10c5a71f625	localstack/localstack	"docker- entrypoint.sh"	2 minutes ago	Up 2 minutes (healthy)	127.0.0.1:4510- 4560->4510- 4560/tcp, 127.0.0.1:4566- >4566/tcp, 5678/tcp
7f0fa023ac4f	3a669f02efff	"python app.py"	7 hours ago	Exited (255) 5 minutes ago	8080/tcp, 0.0.0.0:5002- >5000/tcp
9ff472da8892	3a669f02efff	"python app.py"	7 hours ago	Exited (255) 5 minutes ago	8080/tcp, 0.0.0.0:5001- >5000/tcp

6. Running a PostgreSQL Container

Command:

```
C:\Users\rawat>docker run --name my-postgres -e
POSTGRES_USER=admin -e POSTGRES_PASSWORD=password -e
POSTGRES_DB=mydb -p 5432:5432 -d postgres:15
```

Error Output:

```
docker: Error response from daemon: driver failed
programming external connectivity on endpoint my-postgres
(feae7f0fb87909bde1853a7ddefa49bb518f11250e54304f75109
68f7a88cca1): Bind for 0.0.0.0:5432 failed: port is already allocated.
```

7. Resolving Port Conflict and Running PostgreSQL on a Different Port

Command:

```
C:\Users\rawat>docker run --name my-new-postgres -e
POSTGRES_USER=admin -e POSTGRES_PASSWORD=password -e
POSTGRES_DB=mydb -p 5433:5432 -d postgres:15
```

Output:

```
b2efdca3c6f0af6cf4154fce236f0b66b5efba0f4f9e14972c94b3e0a5afa9de
```

8. Verifying Running Containers

Command:

```
C:\Users\rawat>docker ps
```

Output:

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
b2efdca3c6f0	postgres:15	"docker-entrypoint.s..."	42 seconds ago	Up 41 seconds	0.0.0.0:5433->5432/tcp

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
a10c5a71f625	localstack/localstack	"docker- entrypoint.sh"	3 minutes ago	Up 3 minutes (healthy)	127.0.0.1:4510- 4560->4510- 4560/tcp, 127.0.0.1:4566 >4566/tcp, 5678

## 9. Connecting to PostgreSQL and Performing SQL Operations

Command:

```
C:\Users\rawat>docker exec -it my-  
new-postgres psql -U admin -d mydb
```

Output:

```
psql (15.12 (Debian 15.12-1.pgdg120+1))  
Type "help" for help.
```

Creating a Table and Inserting Data:

```
CREATE TABLE students (  
    id SERIAL PRIMARY KEY,  
    name VARCHAR(100) NOT NULL,  
    email VARCHAR(100) UNIQUE NOT NULL  
);  
INSERT INTO students (name, email) VALUES  
( 'Alice Johnson', 'alice@example.com'),  
( 'Bob Smith', 'bob@example.com'),  
( 'Charlie Brown', 'charlie@example.com');
```

Output:

```
CREATE TABLE  
INSERT 0 3
```

## 10. Performing SQL Queries



## Selecting Data:

```
SELECT * FROM students;
```

## Output:

id	Name	Email
1	Alice Johnson	<a href="mailto:alice@example.com">alice@example.com</a>
2	Bob Smith	<a href="mailto:bob@example.com">bob@example.com</a>
3	Charlie Brown	<a href="mailto:charlie@example.com">charlie@example.com</a>

## Updating Data:

```
UPDATE students SET email = 'bob.smith@example.com'  
WHERE name = 'Bob Smith';
```

## Output:

```
UPDATE 1
```

## Deleting Data:

```
DELETE FROM students WHERE name = 'Charlie Brown';
```

## Output:

```
DELETE 1
```

## Selecting Data with a Condition:

```
SELECT * FROM students WHERE name LIKE 'A%';
```

## Output:

id	Name	Email
1	Alice Johnson	<a href="mailto:alice@example.com">alice@example.com</a>

Exiting PostgreSQL:

\q